

GenAI Research Assistant Hiring Assignment

NL2SQL E-Commerce Query Generator

Objective: Build an end-to-end Natural Language to Query (NL2Query) system that translates natural language questions into safe, executable SQL and Cypher queries over an ecommerce dataset using open-source LLMs, retrieval-augmented generation (RAG), and an **agentic orchestration** framework.

Note: Single purpose LLM will be considered as incomplete assignment

1. Pick any Agentic framework

- a. Phidata (preferred)
- b. Agno
- c. Langchain

2. Database Setup

- a. Use Kaggle E-Commerce Dataset: [Ecommerce Dataset for Analysis](#)
- b. What to deliver:
 - i. Setup a knowledge base of 5-8 tables with Neo4j (graph database)
 - ii. Relationship (can be hand-drawn/exported from neo4j)

3. Vector Database Setup & Prompt Engineering

- a. Model Selection: Choose any model which is freely available
- b. Prompt Engineering (Required):
 - i. Dynamic schema injection (show relevant tables/columns)
 - ii. 10 few-shot examples (NL question → SQL query) stored in vector database, choose any one → Weaviate, Pinecone, FAISS and retrieve top-5 from vector database and dynamically ingest in the prompt
 - iii. Clear instructions for SQL generation
- c. What to deliver:
 - i. Working model integration with code
 - ii. Prompt template with schema injection

- iii. Comparison of 2 different prompt strategies

- d. **Questions to answer:**

- i. Why did you choose this model?
- ii. How does schema context size affect results?

4. Query Validation & Execution

- a. **Validation Pipeline:**

- i. Syntax validation
- ii. Table/column existence - Verify against schema
- iii. Safety checks - Block DROP, DELETE, UPDATE, TRUNCATE
- iv. Query execution - Execute and return results
- v. Self-correction: If query fails, send error back to LLM for retry

- b. **What to deliver:**

- i. Validation module with error handling
- ii. Example of 3-5 validation catches (syntax error, wrong table, etc.)

- c. **Questions to answer:**

- i. What percentage of queries are valid on the first attempt?
- ii. How do you prevent SQL injection?

5. Testing & Evaluation

- a. **Test Suite:** Create **15-20 test queries** across difficulty levels easy, medium and difficult:

- b. **Evaluation Metrics**

- i. **Execution success rate** (% queries that run without errors)
- ii. **Average generation time** (milliseconds)
- iii. **Average execution time** (milliseconds)

- c. **What to deliver:**

- i. Test results table (query, success/fail, error type)
- ii. Performance metrics summary

6. API & Docker Deployment

a. Generate and execute SQL

```
POST /api/nl2sql
{
  "question": "Show total sales by category",
  "execute": true
}
```

b. Get schema info

```
GET /api/schema
```

c. Validate SQL

```
POST /api/validate
{
  "sql_query": "SELECT * FROM products"
}
```

Docker Setup (Optional)

`docker-compose.yml` with 3 services:

```
services:
  database:
    image: postgres:15
    environment:
      POSTGRES_DB: ecommerce
      POSTGRES_USER: user
      POSTGRES_PASSWORD: password
    volumes:
      - ./data/init.sql:/docker-entrypoint-initdb.d/init.sql

  ollama:
    image: ollama/ollama:latest
    volumes:
      - ollama-data:/root/.ollama

  api:
    build: ./api
    ports:
      - "8000:8000"
    environment:
      DATABASE_URL: postgresql://user:password@database:5432/ecommerce
      OLLAMA_URL: http://ollama:11434
```

```
MODEL_NAME: codellama:7b
depends_on:
  - database
  - ollama

volumes:
  ollama-data:
```

What to deliver:

1. Working FastAPI with 3 endpoints
2. Docker setup that runs with docker-compose up
3. API documentation (can use FastAPI auto-docs at /docs)

Deliverables Checklist

1. GitHub Repository

```
nl2sql-ecommerce/
├── data/
│   ├── schema.sql      # Database schema
│   ├── seed_data.sql    # Sample data
│   └── relationship_diagram.png  # relationship diagram
├── nl2sql/
│   ├── generator.py    # NL2SQL generation
│   ├── validator.py    # Query validation
│   └── executor.py     # Query execution
├── api/
│   ├── main.py         # FastAPI app
│   └── Dockerfile
└── tests/
    ├── test_queries.json # 15-20 test cases
    └── test_results.md   # Evaluation results
└── notebooks/
    └── evaluation.py    # Testing & analysis
└── docker-compose.yml
└── requirements.txt
└── .env.example
└── README.md
```

2. README.md

1. Project overview
2. Setup instructions: docker-compose up
3. API usage examples (curl commands)
4. Model used and why
5. Test results summary
6. Known limitations

3. Technical Report (1-2 pages)

File: REPORT.md

Sections:

1. **Database** - Dataset choice, schema description
2. **Model & Prompting** - Model selection, prompt strategy
3. **Validation** - Safety checks implemented
4. **Evaluation Results** - Success rate, examples, failures
5. **Challenges** - Top 3 issues and solutions

4. Demo (Optional but Recommended)

1. 3-5 minute video OR
2. 10 screenshots showing system in action

Submission Instructions

Email with subject: "GenAI RA - NL2SQL Assignment - [Your Name]"

Include:

1. GitHub repository URL (public)
2. Brief summary (3-4 sentences)
3. Setup verification: "Tested docker-compose up on clean machine"
4. Time spent on assignment

Key Questions to Answer

Answer these in your code comments or report:

1. Why is NL2SQL harder than general text generation?
2. Why is RAG-style prompting better than fine-tuning for this task?
3. What percentage of queries succeed on the first attempt?
4. What's the most common failure type?
5. How would you scale this to 50+ tables?

Good luck! Focus on getting a working end-to-end system over perfection.

